

ASYNC: A Cloud Engine with Asynchrony and History for Distributed Machine Learning

**Saeed Soori¹, Bugra Can²,
Mert Gurbuzbalaban², Maryam Mehri
Dehnavi¹**

sasoori@cs.toronto.edu

University of Toronto¹, Rutgers University²



MACHINE LEARNING ON CLOUD

- Machine learning algorithms need to scale on cloud systems.
- Cloud systems are cheap but faulty.
- Nodes can be slow and stall the server.

Finance sciences




Traffic control



Medical sciences



Web related applications



OPTIMIZATION IN MACHINE LEARNING

- In many machine learning application, a model is obtained by applying an **optimization** algorithm on noisy data.
- Recent optimization algorithms must support **asynchrony** and **history** to mitigate the effect of slow workers and noise.
- Existing frameworks:
 - Create static dependency graphs to implement asynchrony¹.
 - Store history using checkpointing or lineage-based methods².

[1] S. Wang, J. Liagouris, R. Nishihara, P. Moritz, U. Misra, A. Tumanov, and I. Stoica, “Lineage stash: Fault tolerance off the critical path,” in Proceedings of Symposium on Operating Systems Principles, SOSP, vol. 19, 2019.


[2] M. Zaharia et al., “Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing,” in Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation. USENIX Association, 2012, pp. 2–2.



OPTIMIZATION IN MACHINE LEARNING


- Problems with current tools:
 - Limited to asynchrony with fixed execution model throughout the algorithm.
 - Overhead of lineage and checkpointing is high.
- Our work:
 - Supports asynchrony by implementing **dynamic dependency graphs**.
 - Recovers history by partial model updates.
- We build our framework on top of Apache Spark¹.

[1] M. Zaharia et al., “Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing,” in Proceedings the 9th USENIX conference on Networked Systems Design and Implementation. USENIX Association, 2012, pp. 2–2.

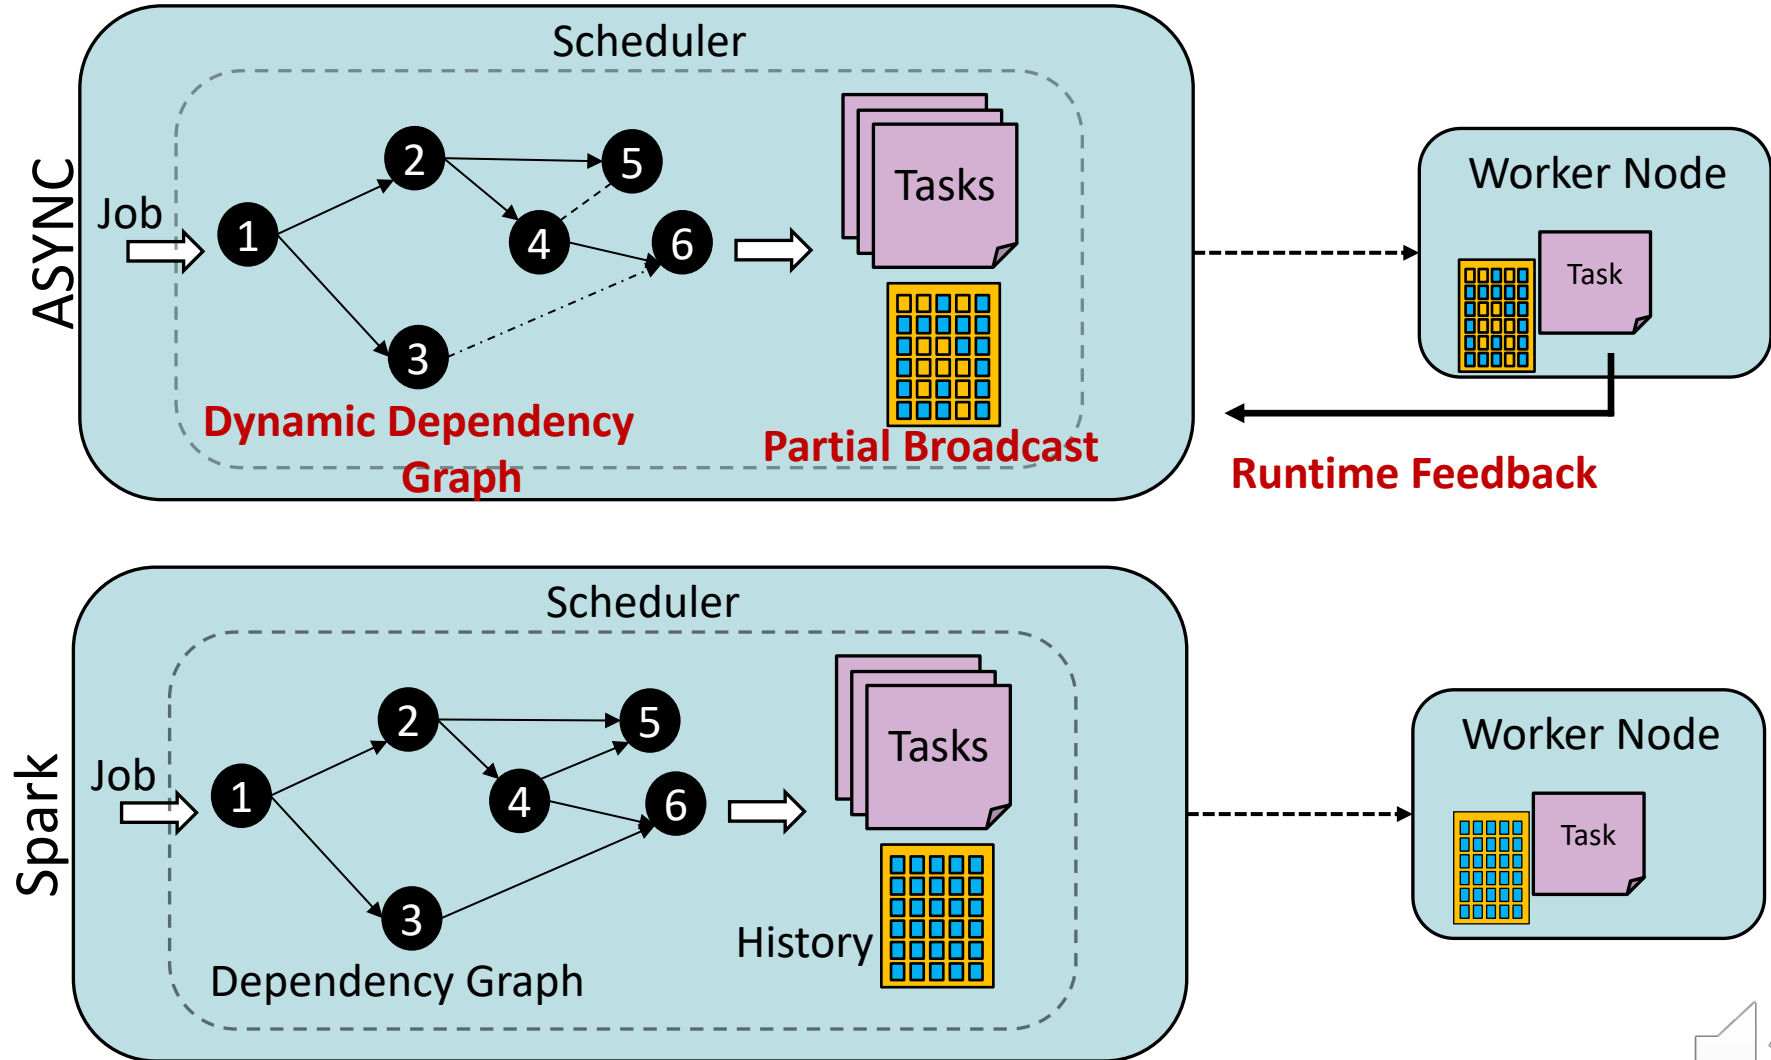


ASYNCR vs SPARK: INTERNALS

- Spark uses a fixed dependency graph with periodic synchronization => **relax the dependencies.**
- It sends the bulky history with tasks => **send partial history and recover the rest.**



ASync vs SPARK: INTERNALS



ASYNC VS SPARK: CODE

- ASYNC uses a programming model written on top of Spark.
- In ASYNC, the user only adds the **blue** lines to convert a synchronous optimization method to an asynchronous variant.

```
Input: points, learning rate  $\alpha$ 
Output: model  $\omega$ 

for i = 1:N
  //broadcast model workers
  broadcast( $\omega$ )
  //compute gradients
   $\nabla f$  = points.map(...).reduce()
  //update model
   $\omega = \omega - \alpha \nabla f$ 
end
```

Spark

```
Input: points, learning rate  $\alpha$ 
Output: model  $\omega$ 

for i = 1:N
  //partial broadcast
  ASYNCbroadcast( $\omega$ )
  //apply custom asynchrony
  points.ASYNCbarrier(...)
   $\nabla f$  = points.map(...).ASYNCreduce()
   $\omega = \omega - \alpha \nabla f$ 
end
```

ASYNC



EXPERIMENTS

- We implement two optimization algorithms in ASYNC and compare to their implementations in Spark:
 - Asynchronous stochastic gradient descent (ASGD)¹: asynchronous without history.
 - ASAGA²: asynchronous with history.
- Datasets are [epsilon](#) and [mnist8m](#) with 400k and 8M samples each³.
- A cluster⁴ of 32 workers with 8 [straggler](#) nodes is used.

[1] L. Bottou, “Stochastic gradient descent tricks,” in Neural networks: Tricks of the trade. Springer, 2012, pp. 421–436.

[2] R. Leblond, F. Pedregosa, and S. Lacoste-Julien, “Asaga: asynchronous parallel saga,” arXiv preprint arXiv:1606.04809, 2016.

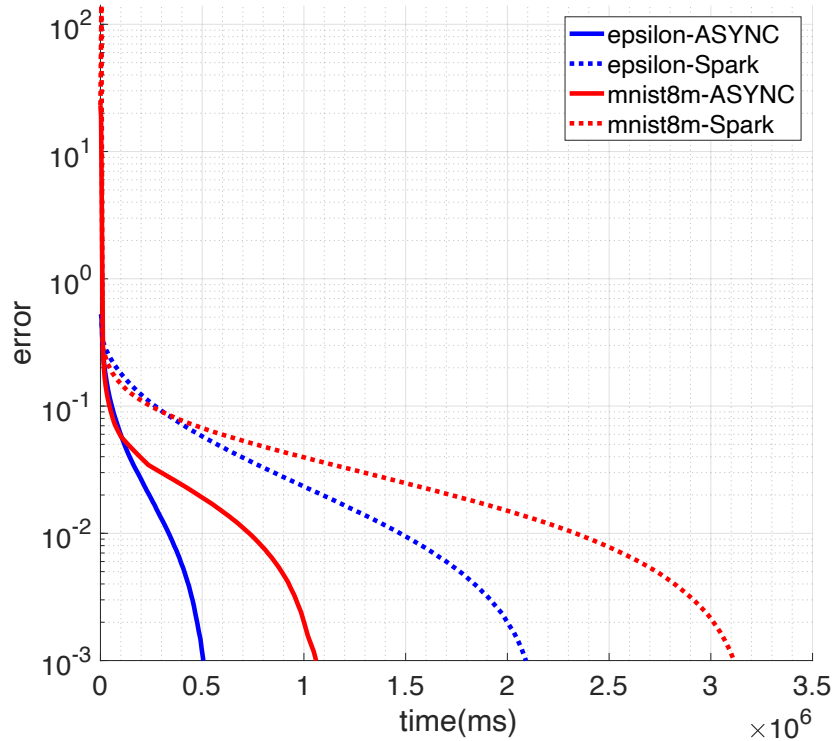
[3] C.-C. Chang and C.-J. Lin, “Libsvm: a library for support vector machines,” ACM transactions on intelligent systems and technology (TIST), vol. 2, no. 3, p. 27, 2011.

[4] J. Towns et al., “Xsede: accelerating scientific discovery,” Computing in Science & Engineering, vol. 16, no. 5, pp. 62–74, 2014.

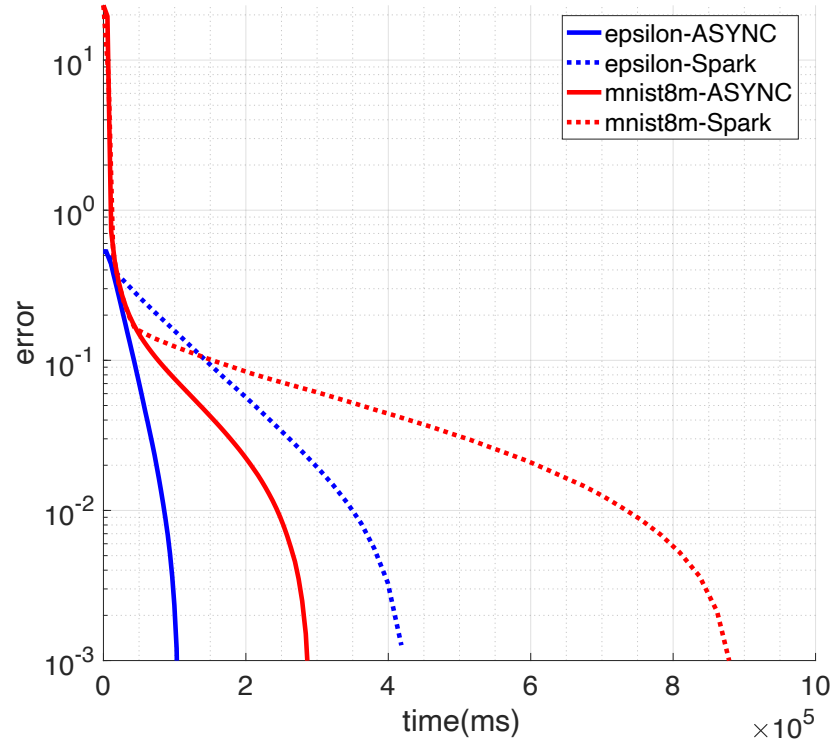


RESULTS

ASGD algorithm



ASAGA algorithm



ASYNC is 4 times faster than Spark.



CONCLUSION

- We developed ASYNC, a framework to support **asynchrony** and **history** for optimization algorithms in machine learning.
- ASYNC outperforms current implementations and achieves **4x** speedup compared to Apache Spark.
- The link to ASYNC:
 - <https://github.com/ASYNCframework/ASYNCframework>
- The link to audio of the presentation:
 - <https://bit.ly/3bh88F6>
 - Email: sasoori@cs.toronto.edu

